


AI KERNEL • V0.1

The open-source foundation for *reliable agents*.

A clean separation of LLM, CPU, and human intelligence.



THE PROBLEM

Current agent systems are *fragile* and *expensive* to operate.

01 The LLM does too much.

Routing, retries, math, state, escalation — all run inside non-deterministic prompts. Output drifts run-to-run.

02 Token waste and runaway cost.


Inference is burned on work the CPU could do for free — validation, schemas, business rules.

03 Locked to AI engineers.

Building or modifying an agent demands prompt engineering expertise. Domain experts can't participate.

04 No audit trail, no oversight.

When something fails, you cannot tell whether the model, the data, or the logic was at fault.




RESULT

High failure rates, exploding inference cost, and an ecosystem only AI engineers can touch.

THE SOLUTION

A declarative DSL that *cleanly separates* the three kinds of intelligence in every agent.



FEATURE 01

Generative UI rendered directly from the DSL schema.

FEATURE 02

Transparent execution graphs: planning → decisions → verification.

FEATURE 03

Full audit log and observability on every run, end-to-end.

WHAT IT LOOKS LIKE

Ten lines of *declarative intent*,
one production-grade pipeline.

claims.ai

AI DSL · v0.1

```

-- Insurance claim intake → extract → draft → flag

SET MODEL gpt-4.1    TEMPERATURE 0    SEED 42

DEFINE claim:
  claimant      TEXT
  policy_number TEXT
  claim_amount  MONEY
  is_disputed   YES/NO
  claim_type    ONE OF [auto, property, health, liability]

FROM      claims/
EXTRACT   claim PROMPT insurance_context EXAMPLES claim_samples
DRAFT     response PROMPT claim_acknowledgment
FLAG WHEN claim_amount OVER 10000
FLAG WHEN claim_type IS liability AND claim_amount OVER 5000
OUTPUT    processed claims.json

```

Same file – readable by the analyst, runnable by the engineer,
diffable in git.

No Python · No YAML · No prompt
engineering.

THE OPPORTUNITY

Unlocking *the next billion*
agent operators.

1B+

ADDRESSABLE OPERATORS

Anyone fluent in SQL or Excel can now build, run,
and operate a production agent.

↓ Tokens

INFERENCE ECONOMICS

Deterministic rules run as code — not prompts.
Dramatically lower per-run cost, dramatically
higher reliability.

OSS

OPEN PRIMITIVES

Parser, compiler, and runtime — open source
foundations the whole ecosystem can build on.

*The missing kernel
for production agents.*



GET IN TOUCH • SHAPE THE PROJECT

David L. Jones

github.com/dljones555/aids1